

```
In [58]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [59]: dataset=pd.read_csv('car data.csv')
```

```
In [60]: dataset.head(2)
```

```
Out[60]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	5.59	27000	Petrol	Dealer	Manual	1
1	sx4	2013	9.54	43000	Diesel	Dealer	Manual	1

```
In [61]: dataset.isnull().sum()
```

```
Out[61]: Car_Name      0
Year      0
Present_Price  0
Kms_Driven  0
Fuel_Type    0
Seller_Type  0
Transmission  0
Owner        0
Selling_Price 0
dtype: int64
```

```
In [62]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null   object
1   Year            301 non-null   int64
2   Present_Price   301 non-null   float64
3   Kms_Driven      301 non-null   int64
4   Fuel_Type       301 non-null   object
5   Seller_Type     301 non-null   object
6   Transmission    301 non-null   object
7   Owner           301 non-null   int64
8   Selling_Price   301 non-null   float64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

## Car\_Name

```
In [63]: from sklearn.preprocessing import LabelEncoder
```

```
In [64]: LE= LabelEncoder()
```

```
In [65]: LE
```

```
Out[65]: LabelEncoder()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [66]: dataset['Car_Name']=LE.fit_transform(dataset['Car_Name'])
```

```
In [67]: dataset.head()
```

```
Out[67]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	90	2014	5.59	27000	Petrol	Dealer	Manual	(
1	93	2013	9.54	43000	Diesel	Dealer	Manual	(
2	68	2017	9.85	6900	Petrol	Dealer	Manual	(
3	96	2011	4.15	5200	Petrol	Dealer	Manual	(
4	92	2014	6.87	42450	Diesel	Dealer	Manual	(

## Fuel\_Type

```
In [68]: dataset['Fuel_Type'].unique()
```

```
Out[68]: array(['Petrol', 'Diesel', 'CNG'], dtype=object)
```

```
In [ ]:
```

```
In [69]: dataset['Fuel_Type']=LE.fit_transform(dataset['Fuel_Type'])
```

```
In [70]: dataset.head(1)
```

```
Out[70]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	90	2014	5.59	27000	2	Dealer	Manual	(

## Seller\_Type

```
In [71]: dataset['Seller_Type']=LE.fit_transform(dataset['Seller_Type'])
```

In [72]: `dataset.head(1)`

Out[72]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owne
0	90	2014	5.59	27000	2	0	Manual	(

## Transmission

In [73]: `dataset['Transmission']=LE.fit_transform(dataset['Transmission'])`

In [74]: `dataset.head(1)`

Out[74]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owne
0	90	2014	5.59	27000	2	0	1	(

In [75]: `x=dataset.iloc[:, :-1]`

In [76]: `x`

Out[76]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Ow
0	90	2014	5.59	27000	2	0	1	
1	93	2013	9.54	43000	1	0	1	
2	68	2017	9.85	6900	2	0	1	
3	96	2011	4.15	5200	2	0	1	
4	92	2014	6.87	42450	1	0	1	
...	...	...	...	...	...	...	...	...
296	69	2016	11.60	33988	1	0	1	
297	66	2015	5.90	60000	2	0	1	
298	69	2009	11.00	87934	2	0	1	
299	69	2017	12.50	9000	1	0	1	
300	66	2016	5.90	5464	2	0	1	

301 rows × 8 columns

In [77]: `y=dataset['Selling_Price']`

In [78]:

y

Out[78]:

```
0      3.35
1      4.75
2      7.25
3      2.85
4      4.60
```

...

```
296     9.50
297     4.00
298     3.35
299    11.50
300     5.30
```

Name: Selling\_Price, Length: 301, dtype: float64

In [79]:

from sklearn.preprocessing import StandardScaler

In [80]:

SS=StandardScaler()

In [81]:

SS

Out[81]:

StandardScaler()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [82]:

x=pd.DataFrame(SS.fit\_transform(x), columns=x.columns)

In [83]:

x.head()

Out[83]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission
0	1.074323	0.128897	-0.236215	-0.256224	0.500183	-0.737285	0.39148
1	1.191828	-0.217514	0.221505	0.155911	-1.852241	-0.737285	0.39148
2	0.212627	1.168129	0.257427	-0.773969	0.500183	-0.737285	0.39148
3	1.309332	-0.910335	-0.403079	-0.817758	0.500183	-0.737285	0.39148
4	1.152659	0.128897	-0.087890	0.141743	-1.852241	-0.737285	0.39148



In [84]:

from sklearn.model\_selection import train\_test\_split


In [85]:

x\_train,x\_test,y\_train,y\_test=train\_test\_split(x,y,test\_size=0.2,random\_sta

```
In [86]: from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
```

```
In [87]: sns.heatmap(data=dataset.corr(),annot=True)
```

```
Out[87]: <Axes: >
```



## LinearRegression()

```
In [88]: LR=LinearRegression()
```

```
In [89]: LR
```

```
Out[89]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [90]: LR.fit(x_train,y_train)
LR.score(x_train,y_train)*100 , LR.score(x_test,y_test)*100
```

```
Out[90]: (88.40630578239454, 84.65539666857805)
```

## Lasso

```
In [91]: LR1= Lasso(alpha=0.05)
LR1.fit(x_train,y_train)
LR1.score(x_train,y_train)*100 , LR1.score(x_test,y_test)*100
```

Out[91]: (88.35433202380113, 84.42023265451037)

## Ridge

```
In [92]: LR2= Ridge(alpha=10)
LR2.fit(x_train,y_train)
LR2.score(x_train,y_train)*100 , LR2.score(x_test,y_test)*100
```

Out[92]: (88.28628537091495, 84.16213595432282)

## ElasticNet

```
In [93]: LR3= ElasticNet(alpha=0.5)
LR3.fit(x_train,y_train)
LR3.score(x_train,y_train)*100 , LR3.score(x_test,y_test)*100
```

Out[93]: (84.00059239671332, 78.3177718663528)

## Decision Tree

```
In [94]: dt=DecisionTreeRegressor(max_depth=13)
dt.fit(x_train,y_train)
dt.score(x_train,y_train)*100 , dt.score(x_test,y_test)*100
```

Out[94]: (99.99983614119861, 95.47372436648206)

```
In [95]: from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [96]: mean_squared_error(y_test,dt.predict(x_test)), mean_absolute_error(y_test,d
```

Out[96]: (1.042654049180328, 0.6585737704918032)

## Random Forest Regressor

```
In [97]: rf=RandomForestRegressor(n_estimators=100)
rf.fit(x_train,y_train)
rf.score(x_train,y_train)*100 , rf.score(x_test,y_test)*100
```

Out[97]: (98.30481566012408, 96.41847578267432)

```
In [98]: mean_squared_error(y_test,rf.predict(x_test)), mean_absolute_error(y_test,rf
```

```
Out[98]: (0.8250250381967219, 0.5829622950819674)
```

## Support Vector regression

```
In [99]: SV=SVR()
SV.fit(x_train,y_train)
SV.score(x_train,y_train)*100 , SV.score(x_test,y_test)*100
```

```
Out[99]: (66.00840380338376, 78.48466914602926)
```

## KNeighborsRegressor

```
In [100]: Knn=KNeighborsRegressor()
Knn.fit(x_train,y_train)
Knn.score(x_train,y_train)*100 , Knn.score(x_test,y_test)*100
```

```
Out[100]: (91.06681012800678, 93.29996797075346)
```

```
In [ ]:
```

```
In [101]: rf.predict([[-1.275759,0.821718,-0.817924,-0.333500,0.500183,1.356327,-2.55
```

C:\Users\aditya\anaconda3\New folder\Lib\site-packages\sklearn\base.py:464:  
UserWarning: X does not have valid feature names, but RandomForestRegressor was fitted with feature names  
warnings.warn(

```
Out[101]: array([0.4543])
```

```
In [102]: y_test
```

```
Out[102]: 177      0.35
289     10.11
228      4.95
198      0.15
60       6.95
...
234      5.50
296      9.50
281      2.10
285      7.40
182      0.30
Name: Selling_Price, Length: 61, dtype: float64
```

```
In [107]: new_data=pd.DataFrame([['ritz',2014,5.59,27000,'Petrol','Dealer','Manual',0
new_data
```

```
Out[107]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner_Type
0	0	2014	5.59	27000	Petrol	Dealer	Manual	0

In [ ]:

In [109]: new\_data['Fuel\_Type']=LE.fit\_transform(new\_data['Fuel\_Type'])

In [110]: new\_data['Seller\_Type']=LE.fit\_transform(new\_data['Seller\_Type'])

In [111]: new\_data['Transmission']=LE.fit\_transform(new\_data['Transmission'])

In [112]: new\_data

Out[112]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner_Type
0		0	2014	5.59	27000	0	0	0

In [113]: new\_data=pd.DataFrame(SS.transform(new\_data), columns=new\_data.columns)

In [114]: new\_data

Out[114]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner_Type
0	-2.4508	0.128897	-0.236215	-0.256224	-4.204665	-0.737285	-2.554408	-0

In [115]: rf.predict(new\_data)

Out[115]: array([3.4745])

In [ ]:

In [ ]: